


Native local targets: 21× faster on CUDA

MATS · NVIDIA L40 · FP64 · 512 tokens · one MLP site in a four-layer language model



All elementwise checks passed; largest absolute discrepancy $9.3\text{e-}13$.

Final input/target transfers included. Model import, compilation and warmup excluded.

Measures supervision capture only—not full training, the new resident mixture, or mechanism discovery. Slurm job 16419.